

Template **this**!

(Divining the value category of ***this**)

Document #: DxxxxR0
Date: 2017-10-24
Audience: Evolution Working Group
Reply-to: Ben Deane, ben at elbeno dot com
Gašper Ažman, gasper dot azman at gmail dot com

Contents

1	Introduction	1
2	Motivation	1
3	Design Considerations	2
4	Proposed Solution	2
5	What does this in a parameter list mean?	3
6	Implications for lambdas	6
7	Impact on the Standard	6
8	Acknowledgments	6
	References	7
A	forward_like	7

1 Introduction

We propose a new mechanism for specifying the value category of an instance of a class which is visible from inside a member function of that class. In other words, a way to tell from within a member function whether one's **this** points to an rvalue or an lvalue, and whether it is **const** or **volatile**.

2 Motivation

The existing mechanism for this (adding a reference, rvalue-reference, **const** or **volatile** qualifier (*cv-ref qualifiers*) suffixes to a member function) suffers from problems.

It is verbose.

A common task in writing a class is providing a "getter function" to access a contained member. Authors that care about performance want to provide getter functions that take advantage of move semantics in the case where the class instance is an rvalue, and so end up writing several of these functions, which only differ in the *cv-ref* qualifiers.

There is no way to write a reference qualifier for a lambda expression.

When writing a lambda expression, it is impossible to know whether it is safe to move from captured-by-copy members of the closure object. This impacts the performance of lambda expressions. One scenario where this is important arises when using lambda expressions as parts of a chain of asynchronous continuations that carry a value through a computation.

There is no way to refer to a lambda expression from within itself.

The rationale in [P0839] (*Recursive lambdas* by Richard Smith) is valid and we'd like to transclude it here.

This paper provides an alternative to the proposed solution in that paper.

TBD: Ask Richard if he'd like to contribute.

3 Design Considerations

In addition to solving the existing problems, desirable properties of a solution are:

- It should work the same way for member functions and lambda expressions.
- It should work like existing practice as much as possible, while adding as little extra syntax as possible.
- It should avoid adding extraneous syntax to function declarations. In a world where every function is fast becoming tagged with attributes, `constexpr`, `noexcept` specifications, and (this proposal notwithstanding) reference qualifiers, readable, noise-free function declarations are to be prized.

TBD: more here.

4 Proposed Solution

We propose the ability to add an optional first parameter to any member function of a class `T`, taking the form `T [const] [volatile] [&|&&] this`.

To facilitate use in generic lambda expressions, this may also be formulated as `auto [const] [volatile] [&|&&] this`.

In all cases, the value category of `this` inside the member function is exactly what the existing parameter rules would already imply. In other words, the *cv-ref qualifiers* that stand after the function signature now explicitly apply to the `this` parameter.

```
1 // Within a class
2
3 struct Person {
4     std::string name;
5
6     // Getter:
7     template <typename U>
8     decltype(auto) GetName(U&& this) {
9         // U can only deduce to Person (const, volatile, &, &&)
10        // because of visibility
11        return std::forward<T>(this).name;
12    }
13
14    // This template will likely instantiate the following functions:
15    // const std::string& GetName() const &;
16    // std::string& GetName() &;
```

```

17     // std::string&& GetName() &&;
18 };
19
20 void foo()
21 {
22     string s = "Hello, world!";
23
24     // Here, the closure object is an rvalue:
25     // the captured s will be correctly moved
26     bar([s] (auto&& this) {
27         return std::forward_like<decltype(this)>(s); // see appendix A
28     })
29 }

```

5 What does **this** in a parameter list mean?

The meaning of the different ways to pass **this** is the same as current general parameter handling.

The entries of this table should be read as if they are inside a class T:

```

1 class T { /* entry */ };

```

In other words, T is *not* a template parameter.

written as	C++17 signature	comments
void f(T this)	currently not available	[value]
void f(T& this)	void f() &	
void f(T&& this)	void f() &&	
void f(T const this)	currently not available	[value]
void f(T const& this)	void f() const&	
void f(T const&& this)	void f() const&&	
void f(T volatile this)	currently not available	[value]
void f(T volatile& this)	void f() volatile&	
void f(T volatile&& this)	void f() volatile&&	
void f(T const volatile this)	currently not available	[value]
void f(T const volatile& this)	void f() const volatile&	
void f(T const volatile&& this)	void f() const volatile&&	

Notes:

- *[value]*: whether passing by value should be allowed is debatable, but seems desired for completeness and parity with inline friend functions.
- The interpretation of **this** in the method body differs, but only one definition for a given signature may be present, eg. one may define at most one of `void f()&`, or `void f(T& this)` or `void f()`, the first and last already being exclusive of one another.

How does templated **this** work?

Using existing deduction rules for template parameters, which will deduce the type of **this** to something in the above table.

What does **this** mean in the body of a member function?

It behaves exactly as a regular parameter declared in the same way.

Constructors

No exceptions to the above rules. If a particular constructor signature is not allowed by the language, it continues to be disallowed. We can already access already-initialized members in initialization lists, which means `this` is already available, even though it hasn't been completely constructed yet.

What about pass-by-value methods?

We think they are a logical extension of the mechanism, and would go a long way towards making methods as powerful as inline friend functions, with the only difference being the call syntax.

One implication of this is that the `this` parameter would be move-constructed in the case where the object is an rvalue, allowing you to treat chained builder methods that return a new object uniformly without having to resort to templates.

Example:

```
1  class string_builder {
2      std::string s;
3
4      operator std::string (string_builder this) {
5          return std::move(s);
6      }
7      string_builder operator*(string_builder this, int n) {
8          assert(n > 0);
9
10         s.reserve(s.size() * n);
11         auto const size = s.size();
12         for (auto i = 0; i < n; ++i) {
13             s.append(s, 0, size);
14         }
15         return this;
16     }
17     string_builder bop(string_builder this) {
18         s.append("bop");
19         return this;
20     }
21 };
22
23 // this is optimally efficient as far as allocations go
24 std::string const x = (string_builder{"asdf"} * 5).bop();
```

Of course, implementing this example with templated `this` methods would have been more efficient due to just having fewer objects, but we got rid of all references in the program!

`this` as a reference

This paper turns `this` into a reference on an opt-in basis, which is in line with the existing guidance that never-null pointers should be references if at all possible, and in this case, it is possible.

We believe there would be no confusion, as in all cases, the value category of `this` is stated plainly in the parameter list, which is on the very same screen.

Teaching also becomes easier, as the meaning of what a “const method” is becomes more obvious to students.

One can always obtain the address of the object by taking the address of the `this`.

Unification with `inline friend` functions

This proposal also makes `this` far less special. In fact, it completely unifies `inline friend` functions and class methods, with the differences being:

- the calling syntax (method vs free function)
- methods can be virtual

Basically, if the first parameter is called `this`, one can parse and instantiate the declaration with exactly the same rules as an `inline friend`, except with a calling convention for methods.

Opt-in uniform call syntax

The interaction of a `this` parameter with `friend` functions raises the possibility of opt-in uniform call syntax. Consider:

```
1 class Foo
2 {
3     friend auto ufcs(Foo this, int x)
4     {
5         // this function can be called two ways
6     }
7 };
```

A reasonable interpretation for defining a member function with `friend` *and* a `this` parameter is that it may be called as either a regular member function, or as a friend function. That is, either syntax could work.

```
1 Foo f;
2 f.ufcs(42); // member function call syntax
3 ufcs(f, 42); // friend (free) function call syntax
```

virtual

We are not entirely sure how `virtual` and value calling conventions would work, so perhaps we can disallow that use case, if the pass-by-value signatures are kept in.

Teachability implications

Using `auto&& this` follows existing patterns for dealing with forwarding references.

Optionally adding "this" as the first parameter fits with many programmers' mental model of the `this` pointer being the first parameter to member functions "under the hood" and is comparable to usage in other languages, e.g. Python and Rust.

It also works as a more obvious way to teach about how `std::bind` and `std::function` with a method pointer work by making the pointer explicit.

ABI implications for `std::function` and related

If references and pointers do not have the same representation for methods, this effectively says "for the purposes of `this`, they do."

6 Implications for lambdas

Generic lambdas, should they take an `auto&& this` parameter, work according to existing rewriting rules: the `auto&& this` is turned into a "forwarding reference" and deduced as if it were inside a `template <typename T> auto operator()(T&& this) { ... }`.

Interplays with capturing `[this]` and `[*this]`

As `this` is passed as an explicit parameter, it shadows members of the closure. That said, it's not possible to refer to the members of the closure using the `this` pointer of the lambda, since the closure has no defined layout and the members referenced may not even be inserted into the closure, if there even is one.

TBD: does init-capture obviate all need for `*this`?

Do we allow `this` in lambdas that decay to a function pointer?

If the lambda would otherwise decay to a function pointer, `&this` shall have the value of that function pointer.

Does `this` allow recursion in lambdas?

Yup. You're allowed to call `this(...)`.

Expressions allowed for `this` in lambdas:

```
1  this(...);    // call with appropriate signature
2  decltype(this); // evaluates to the type of the lambda with the appropriate
3                  // cv-ref qualifiers
4  &this;        // the address of either the closure object or function pointer
5  std::move(this) // you're allowed to move yourself into an algorithm...
6  /* ... and all other things you're allowed to do with the lambda itself. */
```

TBD: I don't think we should say that you can now refer to the closure object. What you can do is deduce the value category of the closure object, and access its members in a way appropriate to it.

TBD: Do we believe this?

7 Impact on the Standard

TBD: A bunch of stuff in section 8.1.5 [expr.prim.lambda].

TBD: A bunch of stuff in that `this` can appear as the first method parameter.

8 Acknowledgments

Thanks to the following for their help and guidance:

- Louis Dionne
- Marshall Clow

References

[P0839] Richard Smith, *Recursive lambdas*
<http://wg21.link/p0839r0>

[P0018r3] *Capturing *this by value* <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2016/p0018r3.html>

[P0637r0] *Capture *this With Initializer* <http://www.open-std.org/jtc1/sc22/wg21/docs/papers/2017/p0637r0.html>

A forward_like

We propose a new library facility, `forward_like`, that acts as `std::forward<T>(t).member` does for members, with the syntax `forward_like<T>(t.member)`. This is done because in the case of lambdas, the closure members are not actually addressable using `this.member`.

The proposed semantics are essentially the same as `std::forward`. If `From` and `To` are the same type, `forward_like` and `forward` act exactly the same.

Proposed definition:

```
33 template <typename Like, typename T>
34 constexpr decltype(auto) forward_like(T&& t) noexcept
35 {
36     // first, get 't' back into the value category it was passed in
37     // then, forward it as if its value category was 'Like's.
38     // This prohibits rvalue -> lvalue conversions.
39     return std::forward<like_t<Like, T>>(std::forward<T>(t));
40 }
```

To do that, we require another facility in the standard library, `like`:

```
9  template <typename From, typename To>
10  class like {
11      template <bool Condition, template <typename> class Function, typename T>
12      using apply_if = std::conditional_t<Condition, Function<T>, T>;
13      using base = std::remove_cv_t<std::remove_reference_t<To>>;
14      using base_from = std::remove_reference_t<From>;
15
16      static constexpr bool rv = std::is_rvalue_reference_v<From>;
17      static constexpr bool lv = std::is_lvalue_reference_v<From>;
18      static constexpr bool c = std::is_const_v<base_from>;
19      static constexpr bool v = std::is_volatile_v<base_from>;
20
21  public:
22      using type = apply_if<lv, std::add_lvalue_reference_t,
23                          apply_if<rv, std::add_rvalue_reference_t,
24                          apply_if<c, std::add_const_t, apply_if<v, std::add_volatile_t, base>>>>;
25  };
26
27  template <typename From, typename To>
28  using like_t = typename like<From, To>::type;
```

It merely copies the cv-ref qualifiers from `From` to `To`.

The entire listing of the code with all the tests is available at https://github.com/atomgalaxy/isocpp-template-this/blob/master/forward_like.cpp.